

Applied Cryptography

Ivan Valentini

Telegram: @IvanV1337

Github: https://github.com/IvanValentini/Applied_Cryptography

January 24, 2023

Contents

1	Introduction to Cryptography	3
1.1	Introduction	3
1.2	Attacks	4
1.2.1	Ciphertext-only attackers (COA)	4
1.2.2	Known-plaintext attackers (KPA)	5
1.2.3	Chosen-plaintext attackers (CPA)	5
1.2.4	Chosen-ciphertext attackers (CCA)	6
1.3	Shannon Theorem	6
1.3.1	Consequences of Shannon Theorem	7
1.4	Vernam Cipher - One Time Pad	7
1.5	From perfect to ideal	8
1.5.1	Practical secure ciphers	8
2	Stream Ciphers	10
2.1	The problem with Vernam cipher	10
2.2	Symmetric encryption	10
2.2.1	Stream ciphers	11
2.2.2	Vigenère cipher	12
2.3	Examples of symmetric ciphers	12
2.3.1	Examples of stream ciphers	12
2.4	Implementation of stream ciphers	13
2.5	Warm Up	14
2.6	Linear feedback shift register (LFSR)	14
2.7	LFSR and finite fields	14
2.8	LFSR and linear algebra	18
2.8.1	DVD encryption	18
2.9	A5 family of ciphers	19
2.9.1	GSM	19
2.9.2	A5/0	20
2.9.3	A5/1	20
2.9.4	A5/2	21
2.9.5	A5/3	21
2.10	Remarks on the A5 Family	21
2.11	Practicality of attacking GSM communication	22
2.12	Bluetooth	22
2.12.1	E0	22

2.13 RC4	23
2.14 RC4 in WEP	25
2.14.1 WIFI Protocol	25
2.14.2 WIFI Authentication	25
2.14.3 Security of WPA2	25

Chapter 1

Introduction to Cryptography

1.1 Introduction

Cryptography refers to hidden writing. Its goal is to enable a secure communication between two users (Alice and Bob), and making it impossible for an eavesdropper to understand the information being exchanged.

By channel we mean any physical or logical medium of communication from one user to another. A channel becomes secure when the information exchanged over it cannot be overheard or tampered with by eavesdroppers. By default a channel is considered insecure, so the intent of cryptography is to make secure, an insecure channel.

To transmit data in a secure way Alice and Bob rather than transmitting the message in a plain form they first covert it to a disguised form. Formally these are called plaintext (\mathcal{P}) and ciphertext (\mathcal{C}). The idea is to transform a plaintext into a ciphertext, so that Alice sends the latter to Bob, and Bob is able to reconstruct the plaintext from the received ciphertext while this is very difficult (almost impossible) for Eve.

In practice there is a pair of functions:

- $enc : \mathcal{P} \rightarrow \mathcal{C}$
- $dec : \mathcal{C} \rightarrow \mathcal{P}$

Such that $dec(enc(m)) = m$ for every $m \in \mathcal{P}$. For this to work Alice and Bob need to agree on what encryption and decryption scheme to use without disclosing it to Eve. Eve will only be able to observe ciphertexts, but notice that if the sets of plaintexts and ciphertexts are too small, then




Figure 1.1

Eve can try all the plaintext-ciphertext pairs (exhaustive search) or even if the sets of plaintexts and ciphertexts are large enough to make exhaustive search impractical, encryption and decryption can be defined in obvious way to allow Eve to easily reconstruct them (guessing). There is a problem with this formalization: these two functions, must be kept secret. These function can be used to create a secure channel, but how do you share these function? We need a secure channel, but if we had a secure channel, why not use it in the first place?

Since otherwise we would have to define an encryption and decryption functions for each pair of people that want to communicate securely we introduce the concept of a cryptographic key. We denote the set of (cryptographic) keys with \mathcal{K} , and consider a map $\varphi : \mathcal{P} \times \mathcal{K} \rightarrow \mathcal{C}$ such that for every key $k \in \mathcal{K}$ the function $\varphi(\cdot, k) : \mathcal{P} \rightarrow \mathcal{C}$ is an encryption function.

There are some key differences with respect to the old functions:

- The definition of φ (encryption algorithm) can be very complex but it can be public (i.e. known to anyone) and Alice and Bob can agree on it over an insecure channel. Before the encryption algorithm had to be private.
- The only component that must be kept secret (and thus exchanged over a secure channel) is the cryptographic key k , that defines the encryption function to use

The Kerckhoffs principle states that a cryptosystem should be secure even if everything about the system, except the key, is public knowledge. The advantages to only exchanging cryptographic keys rather than ciphers are that it is easier to keep secret k than φ , and if the key is discovered it is sufficient to choose another key. The main disadvantage is that the attacker only needs to find the key to break the system.

1.2 Attacks

Let's now consider some useful attack models expressed in terms of what the attacker can observe and what queries they can make to the cipher. A query for our purposes is the operation that sends an input value to some function and gets the output in return, without exposing the details of that function. An encryption query, for example, takes a plaintext and returns a corresponding ciphertext, without revealing the secret key. We call these models black-box models, because the attacker only sees what goes in and out of the cipher. There are several different black-box attack models. Note that the higher the attacker skills and complexity of the attack the less (computational) effort the attacker needs to put to break the system.

1.2.1 Ciphertext-only attackers (COA)

Ciphertext-only attackers (COA), or known-ciphertext attackers, observe ciphertexts but don't know the associated plaintexts, and don't know how the plaintexts were selected. Attackers in the COA model are passive and can't perform encryption or decryption queries. The task of the attacker is very difficult and a lot of computational power is required to mount such an attack, this is because the attacker needs to check for every possible key if the decrypted ciphertext is meaningful 1.2.

In some cases the attacker only needs to know the probability distribution of the plaintexts, it could obtain a lot of information merely by observing some ciphertexts. This holds under the assumption that all plaintexts are encrypted with the same cipher and the same key. The method may be difficult (if possible at all) to apply to short messages or messages that contain words with many occurrences of letters with low frequencies. More information can be found here.




Figure 1.2




Figure 1.3

1.2.2 Known-plaintext attackers (KPA)

Known-plaintext attackers (KPA) observe ciphertexts and do know the associated plaintexts. Attackers in the KPA model thus get a list of plaintext–ciphertext pairs, where plaintexts are assumed to be randomly selected. Again, KPA is a passive attacker model, thus the attacker can not choose a plaintext to encrypt.

Essentially the attacker will attempt to do what is known as a key recovery attack. But recovering the key might be a difficult problem to solve, so Eve might try to discover a functionally equivalent algorithm for encryption and decryption, or else design cryptographic algorithm that, even without knowing the key k , produces the same result as those of the cipher with the key k . This kind of attacks are known as Global Deduction/Reconstruction attacks 1.3.

1.2.3 Chosen-plaintext attackers (CPA)

Chosen-plaintext attackers (CPA) can perform encryption queries for plaintexts of their choice and observe the resulting ciphertexts. This model captures situations where attackers can choose all or part of the plaintexts that are encrypted and then get to see the ciphertexts. Unlike COA or KPA, which are passive models, CPA are active attackers, because they influence the encryption processes rather than passively eavesdropping 1.4.

With this the attacker may try to guess previously unknown plaintext-ciphertext pairs.

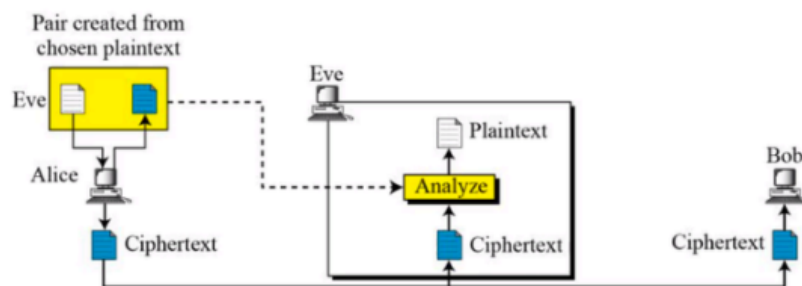


Figure 1.4

1.2.4 Chosen-ciphertext attackers (CCA)

Chosen-ciphertext attackers (CCA)¹ can both encrypt and decrypt; that is, they get to perform encryption queries and decryption queries. The CCA model may sound ludicrous at first—if you can decrypt, what else do you need?—but like the CPA model, it aims to represent situations where attackers can have some influence on the ciphertext and later get access to the plaintext. Moreover, decrypting something is not always enough to break a system.

1.3 Shannon Theorem

A cipher is broken if a method of determining the plaintext from the ciphertext is found without being legitimately given the decryption key. Any cryptosystem can be broken by an exhaustive key search. There exist ciphers that cannot be broken that are called perfect, and even exhaustive key search is of limited use for these ciphers. A cipher is perfect if, after seeing the ciphertext, an attacker gets no extra information about the plaintext other than what was known before the ciphertext was observed. The attacker might know the kind of content hidden but not the content itself, the point is that the knowledge of the attacker about the plaintext is not increased after intercepting the ciphertext. Let:

- \mathcal{K} be the set of keys
- \mathcal{M} be the set of messages
- \mathcal{C} be the set of ciphertexts
- φ is a cipher such that $\forall k \in \mathcal{K}$

$$\varphi_k : \mathcal{P} \rightarrow \mathcal{C}$$

is an encryption function. Furthermore φ_k is invertible: there exists $\varphi_k^{-1}(\varphi_k(m)) = m$. We will give a notion of probabilities:

- if $m \in \mathcal{P}$, we use $P(m)$ to indicate the a priori probability that m occurs.
- if $k \in \mathcal{K}$, $P(k)$, indicates the probability that k is the chosen key

¹Not in the slides, maybe not required to know this

- if $c \in \mathcal{C}$, $P(c)$ is the probability that c is the transmitted ciphertext
- $P(\mathcal{P})$ is the probability distribution of the plaintexts
- $P(\mathcal{C})$ is the probability distribution of the ciphertexts
- $P(\mathcal{K})$ is the probability distribution of the keys

We can assume $P(m) > 0$ because otherwise m never occurs and we can remove it from \mathcal{P} . In the same way we assume $P(k) > 0$ and $P(c) > 0$. The two probability distribution $P(\mathcal{P})$ and $P(\mathcal{K})$ determine the probability distribution $P(\mathcal{C})$, because only one ciphertext can be obtained using the encryption algorithm φ with a given plaintext and key. Clever boy assumption states that plaintext and the keys are independent. So we do not select a particular key for a particular plaintext, for any plaintext we can choose any key. A cipher is called perfect if:

$$\forall m \in \mathcal{P} \wedge \forall c \in \mathcal{C} \implies P(m) = P(m | c)$$

So the probability of having the plaintext m is equal to the conditional probability of observing m over the channel once you have seen the ciphertext c . The conditional probability $P(x | y)$ denotes the probability that x occurs given that y occurred. We say that x and y are independent if $P(x \wedge y) = P(x)P(y)$. The concept of independence formalises the perception that past events do not influence the outcome of future ones or provide any information about them. In other words, if two events are independent, the order in which they occur is of no importance. So seeing the ciphertext c does not add any knowledge about the plaintext message sent, or any future message exchanged.

Shannon theorem: let us fix an integer n . Assume that: keys, plaintext, ciphertext have all n bits. Assume also that the plaintexts and the keys are independent (Clever boy assumption) and any n -bit string may be either a key or a plaintext. Then φ is a perfect cipher if and only if both the following conditions hold:

1. the keys are perfectly random
2. for any pair (m, c) of plaintext-ciphertexts, there is one and only one key k such that $c = \varphi(m, k)$

1.3.1 Consequences of Shannon Theorem

If φ is a perfect cipher, then all ciphertexts have the same probability to be received. Even if Eve knows the probability distribution of the plaintexts, she observes a perfectly random cipher. Hence Eve cannot recover any information on the sent message from the intercepted ciphertext. But we made a very big assumption: this is true only if Eve can intercept **one** ciphertext. If Eve intercepts more ciphertexts encrypted with the same key, then the Shannon theorem no longer guarantees perfect secrecy. So in practical sense, a perfect cipher is not unbreakable, so a perfect cipher is not necessarily ideal, because every time we need to use a different key (the key space has to be as large as the message space) also the key has to be as large as a message transferred. Additionally we did not guarantee any sort of integrity!

1.4 Vernam Cipher - One Time Pad

The One Time Pad is an example of a perfect cipher. We work under the same assumptions of the Shannon Theorem. In the Vernam cipher we define encryption by means of the bitwise XOR




Figure 1.5

operation 1.5:

$$c = \varphi_k(m) = k \oplus m$$

Only if the key used is randomly chosen, this cipher is perfect. But there are a few problems: how do we choose the key randomly, how is the key shared, and the fact length of the key used that has to be as long as the message shared (the storage needed doubles) and having to use a key just once. Let's say that Eve known a single plaintext-ciphertext pair (m_1, c_1) , the key can be easily recovered as $m_1 \oplus c_1 = m_1 \oplus (m_1 \oplus k) = (m_1 \oplus m_1) \oplus k = 0 \oplus k = k$. So any further messages exchanged with this key can be recovered by the attacker.

1.5 From perfect to ideal

We want the ability to encrypt a long message (e.g. a file of several megabytes) using a short key (e.g. a few hundred bits). We do not consider all possible adversaries, but only computationally feasible adversaries, that is, "real world" adversaries that must perform their calculations on real computers using a reasonable amount of time and memory. This leads to a weaker definition of security called semantic security. Furthermore, for our purposes a cipher is to be considered secure long as they do not leak any useful information about an encrypted message to an adversary other than the length of the message. Since the focus is on the "practical", instead of the "mathematically possible", one shall also insist that the encryption and decryption functions are themselves efficient algorithms and not just arbitrary functions. So we'll study ciphers are regarded as secure in practice because the known theoretical attacks take too much time to conduct. In other words, to implement such theoretical attacks requires resources which are unrealistic for any attacker.

1.5.1 Practical secure ciphers

Characterizing the notion of practical security is not an easy task as it must consider several different aspects including:

- Cover time: the time window in which a plaintext must be kept secret. This suggests that no attack on the cipher can be conducted in less than the cover time and implies that an exhaustive key search takes longer than the cover time. Evaluation should be repeated if a new attack is discovered or other parameters are changed such as the available computation power.
- Computational complexity: what computational processes are involved in known attacks on the cryptosystem how much time it takes to conduct these processes. measuring the time taken to perform the processes requires a way of measuring the time it takes to run a process, or a

function and this is expressed as a function of the size of the input that expresses the number of elementary operations performed, known as the Big O notation.

Most modern ciphers are regarded as secure in practice because the known theoretical attacks take too much time to conduct. Exhaustive key search has complexity 2^n , assuming a computer can attempt a million keys per second, an exhaustive search in a 30-bit key space will be covered in about a thousand seconds. Establishing the complexity of any known attacks is important and useful, but brings no guarantees of practical security and this is because there are undiscovered theoretical attacks, problems with the implementation, key management issues and so on. How can we be sure an attacker will require a large amount of work to break a non-perfect system with every method? An alternative, it is to show that breaking the cipher can be reconducted to a computationally difficult problem. This is typically the approach used to show the security of public key ciphers.

Chapter 2

Stream Ciphers

2.1 The problem with Vernam cipher

There are few issues: generating a truly random key as long as the message, find a secure channel for transportation of the key to the message recipient and do this for every single message to be exchanged. In summary, the problem becomes to securely transfer large quantities of secure keys. There are two approaches that are seen as an improvement to Vernam cipher: stream ciphers and block ciphers.

2.2 Symmetric encryption

In stream ciphers we take a seed (a small vector of a few random bits) that must be kept secret, then build a keystream (a very long sequence of pseudorandom bits), finally xor the keystream with the plaintext bitwise to calculate the ciphertext. The problems are how are we going to generate a perfectly random keys of arbitrary size and how can we replicate the random streams of bytes for decryption. In block ciphers we use the same key multiple times in a way that does not compromise the cipher. The problems are how can we reuse multiple times the same key without enabling an attacker to perform cipher-text only attacks and how can we avoid attackers to exploit the block structure.

These two ciphers are known as symmetric encryption algorithms, where stream ciphers perform operations in a way such that the plaintext is processed one bit at a time, and the algorithm selects one bit of plaintext, performs a series of operations on it, and then outputs one bit of ciphertext, block ciphers perform operations in a way such that the plaintext is processed in blocks (groups) of bits at a time, and the algorithm selects a block of plaintext bits (typically 64 bits), performs a series of operations on them, and then outputs a block of ciphertext bits. Notice that a stream cipher can be seen as a block cipher with blocksize set to 1 bit, but there are also stream ciphers that process data in bytes, and hence could be regarded as block ciphers with a block size of 8, as a rule of thumb if the blocksize is less than 64 bits we talk about stream ciphers otherwise we talk about block ciphers.




Figure 2.1: Stream cipher

2.2.1 Stream ciphers

The real work in designing a good stream cipher goes into designing the keystream generator. Keystream generators produce output which appears to be randomly generated, but is actually not randomly generated, these are referred to as pseudorandom generators. In many cases, stream ciphers combine the keystream with the plaintext in more complex ways than a simple bitwise xor operation.

For a stream cipher to be good, Eve should not be able to: recover the seed by making the set of possible seeds so large that an exhaustive search is very hard in practice and also predict the rest of the keystream by eliminating any patterns from the keystream. More precisely: choose a short l -bit (much smaller than the length L of the plaintext to be encrypted) seed s as the encryption key and stretch the seed into a longer L -bit string (the key) that is used to mask the message and decrypt the ciphertext. The seed s is stretched using some efficient, deterministic algorithm G that maps l -bit strings (seeds) to L -bit strings (keys). Formally:

- Encryption: $G(s) \oplus m$ for any seed s (of size l) and plaintext m (of size L)
- Decryption: $G(s) \oplus c$ for any seed s (of size l) and ciphertext c (of size L) where G is called a pseudo-random number generator

Notice that if $l < L$, then by Shannon's Theorem, stream ciphers cannot be perfect, however, if G has certain properties, then stream ciphers are secure in practice. Suppose s is a random l -bit string and r is a random L -bit string, if Eve cannot effectively tell the difference between $G(s)$ and r , then it should not be able to tell the difference between stream ciphers and one-time pad. Since the one-time pad cipher is secure, so should be the stream cipher.

An algorithm that is used to distinguish a pseudo-random string $G(s)$ from a truly random string r is called a statistical test. How might one go about designing an effective statistical test? One basic approach is the following: given an L -bit string, calculate some statistic, and then see if this statistic differs greatly from what one would expect if the string were truly random. For example, a very simple statistic that is easy to compute is the number k of 1's appearing in the string. For a truly random string, we would expect $k \approx L/2$. If the PRG G had some bias towards either 0-bits or 1-bits, we could effectively detect this with a statistical test that, say, outputs 1 if $|k - 0.5L| < 0.01L$, and otherwise outputs 0. This statistical test would be quite effective if the PRG G did indeed have some significant bias towards either 0 or 1.

A stream-cipher is well equipped to encrypt a single message from Alice to Bob. If two messages are encrypted with the same key there may be problems. As an example consider the case in which Alice and Bob want to exchange messages m_1 and m_2 , let $c_1 = m_1 \oplus G(s)$ and $c_2 = m_2 \oplus G(s)$. If Eve is able to intercept both ciphertexts, then it is able to calculate $c_1 \oplus c_2 = (m_1 \oplus G(s)) \oplus (m_2 \oplus G(s)) = (m_1 \oplus m_2) \oplus (G(s) \oplus G(s)) = m_1 \oplus m_2$, and as english text contains enough redundancy that given $m_1 \oplus m_2$, Eve can recover both m_1 and m_2 in the clear by using frequency analysis (given that both are sufficiently long). For this reason a stream cipher key should never be used to encrypt more than one message.

Stream-ciphers are said to be malleable since an attacker can cause predictable changes to the plaintext, this is because an attacker can intercept ciphertext c and forwarding $c' = c \oplus d$, effectively the receiver will get $m' = c' \oplus G(s) = (c \oplus d) \oplus G(s) = (c \oplus G(s)) \oplus d = m \oplus d$. So again stream-ciphers do not provide integrity. Regarding key management, stream ciphers do not require the key to be as long as the message encrypted, like the one time pad does. The key used to generate the keystream (and that must be distributed) is much shorter. In one-time pad the key has to be truly randomly generated, which involves costly generation techniques. The keystream in a stream cipher is pseudorandom and thus is much cheaper to generate. A keystream generator is a deterministic process since every time the same seed is input into the keystream generator, it will result in the same keystream being output. If we reuse a seed to produce the same keystream and then encrypt two plaintexts using the same portion of the keystream, then, just as in a one-time pad, the xor between the two ciphertexts will tell us the difference between the two corresponding plaintexts. We can avoid this problem by, e.g., making the keystream dependent on time varying data or generating the ciphertext with more complex operations than a xor.

In summary stream-ciphers do not give rise to error propagation as each bit in the ciphertext depends on just one bit in the plaintext and 1 bit transmission errors will result in 1 bit error in the plaintext, also they are very fast making them ideal for real time applications (e.g. mobile communication services) and easy to implement in hardware and don't require large memory capabilities. Since stream ciphers process data bitwise, it is crucial that sender and receiver keep their keystreams in perfect synchronization and 1 bit data loss may have catastrophic consequences as decryptions are performed on the wrong bits after the receiver is out of sync of the sender and re-synchronization mechanisms must be put in place to avoid these problems.

2.2.2 Vigenère cipher

It can be seen as a variant of Vernam cipher whereby the key is a sequence of bits of fixed length. The key, and plaintext are a string of bytes, to encrypt: XOR each character in the plaintext with the next character of the key and wrap around in the key as needed.

Vigenere cipher can be attacked by first determining the key length and determining each byte of the key by using frequency analysis.

2.3 Examples of symmetric ciphers

2.3.1 Examples of stream ciphers

- RC4: Simple and fast stream cipher with a relatively low level of security, probably the most widely implemented stream cipher in software and widely used in SSL/TLS, WEP, and Microsoft Office

Type	Cipher	Applications
Stream	A5/1, A5/2, A5/3	Phone (GSM)
	RC4	Internet
	E0	Bluetooth
Block	DES	(old)
	3DES	Smart cards
	AES	Everywhere
	PRESENT	sensor networks

Figure 2.2: Symmetric ciphers

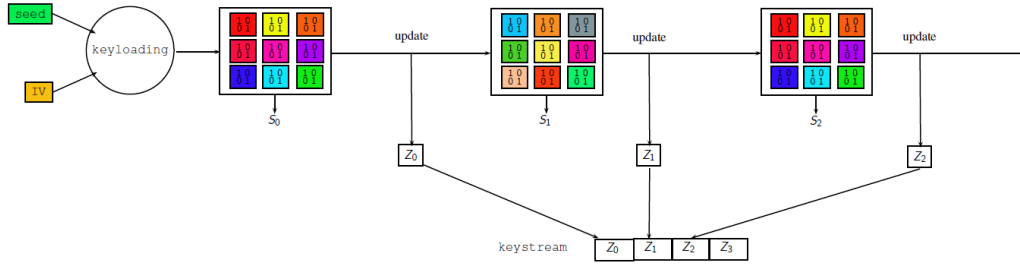


Figure 2.3: Stream ciphers internals

- A5/1: One of the stream cipher algorithms used in GSM to secure the communication channel over the air from a mobile phone to the nearest base station
- E0: The stream cipher algorithm used to encrypt Bluetooth communications

2.4 Implementation of stream ciphers

First we want to tackle the problem of generating a randomly a key (known as keystream) that is as long as possible. Keystream generators should be fast (as in computable in polynomial time as function of number l of bits in the seed) and be secure, so intuitively, a string of L bits produced by a keystream generator should look random. (i.e. it should be impossible in a polynomial amount of time in l to distinguish between a truly random bit string of length L and a string of the same length returned by the keystream generator). More importantly, a keystream generator, given the same seed it should produce the same sequence! The main components are 2.3:

- States: vector of bits organized in registers (S_0, S_1, \dots)
- Update function: function mapping a state to the next state (clock function)
- Output function: function extracting a bit from a state. Concatenating all bits returned by this function, it is possible to obtain the keystream
- Key loading: function that takes the seed (secret) and a (public) initialization vector (IV) to compute the initial state for the update function. Each IV should be used only once.

2.5 Warm Up

The first output bits strongly depend on the initial state. To avoid potential problems, it is customary to run a warm-up phase before starting encryption. This preliminary phase consists in applying the update function several times without outputting any bits of keystream and it is a highly recommended security best practice. Given any initial state, the states are periodic, since they are in a finite number and at some point we will obtain again one of the previous states. The keystream is also periodic and this is impossible to avoid. The smallest number i such that $update(\dots(update(S))) = S$ is called the period of the keystream (it depends on the initial state), and as a requirement is that the period of the keystream shall be quite large, regardless of the initial state, and can be achieved by a suitable design of the update function. As an example let's take an un update function as follows $f : (x, y, z) \Rightarrow (y + z, x, y)$. One can see that the repeated application of f to the initial state $(1, 0, 1)$ will yield $(1, 0, 1) \Rightarrow (1, 1, 0) \Rightarrow (1, 1, 1) \Rightarrow (0, 1, 1) \Rightarrow (0, 0, 1) \Rightarrow (1, 0, 0) \Rightarrow (0, 1, 0) \Rightarrow (1, 0, 1)$ with a period of 7. We are using linear functions because they are easy to implement and compute. Later we will see that linear functions (alone) should never be used.

2.6 Linear feedback shift register (LFSR)

A linear feedback shift register of length n is a shift register composed by n bits. At any clock the following operations are executed:

- the last bit on the right is output and forms part of the keystream
- the other bits in the register are shifted to the right by one position
- the XOR of some bits of the register are put in the first position on the left.

A shift register is a type of digital circuit using a cascade of flip flops (device which stores a single bit of data) where the output of one flip-flop is connected to the input of the next. Flip-flops share a single clock signal, which causes the data stored in the system to shift from one location to the next. A linear-feedback shift register (LFSR) 2.4 is a shift register whose input bit is a linear function of its previous state. The most commonly used linear function of single bits is xor. An LFSR is usually a shift register whose input bit is driven by the xor of some bits of the overall shift register value, and the initial value of the LFSR is called the seed. Since the operation of the register is deterministic, the stream of values produced by the register is completely determined by its current (or previous) state.

Since the register has a finite number of possible states ($2^n - 1$), it must eventually cycle. An LFSR with a well-chosen feedback function can produce a sequence of bits that appears random and has a very long cycle. Typically the linear feedback function has the following form: $c_1q_1 \oplus c_2q_2 \oplus \dots \oplus c_nq_n$ where \oplus denotes the XOR operation. The non null c_i are called taps. Given a seed s , the period of s is the number of steps the LFSR takes to return to s . The period of the LFSR is the maximum period achieved for any seed. For any number n of taps there exists a maximal LFSR. Of the $2^n - 1$ possible LFSRs, which taps correspond to maximal LFSRs? To answer this question we can use finite fields.

2.7 LFSR and finite fields

Note: This is an informal introduction aiming to convey ideas with no attempt to mathematical rigour.




Figure 2.4: Linear feedback shift register

Consider addition modulo a certain number N , this is an instance of the notion of group, namely:

- a set closed under a binary operation (addition modulo N)
- the operation is associative (addition modulo N is so)
- there is an identity (0)
- each element has an inverse (e.g., 1 is the inverse of 11 when considering addition modulo $N=12$ as adding them gives 0)

The order of a group is the number of elements in its set (in the case of addition modulo a number N , the order of the group is N). The order of an element a in a group is the number of times one needs to apply the operation to a to produce the identity (example, the element 4 in the group modulo 12 has order 3 since $4 + 4 + 4 = 0$). Consider now the multiplication modulo some number N . It is interesting to consider the order of its various elements that may be regarded to form cycles, containing the elements obtained by repeatedly applying the operation to the initial element up to its order, covering all (in case the order of the element is N) or some of the N elements:

- 3 has order 3 since $3 * 3 \bmod 13 = 9$, $9 * 3 \bmod 13 = 1$ and $1 * 3 \bmod 13 = 3$, the cycle generated is $\{3, 9, 1\}$
- 2 has order 12 since $2 * 2 \bmod 13 = 4$, $4 * 2 \bmod 13 = 8$, $8 * 2 \bmod 13 = 3$, $3 * 2 \bmod 13 = 6$, and so on, the cycle generated is $\{2, 4, 8, 3, 6, 12, 11, 9, 5, 10, 7, 1\}$

It is not difficult to see that there are other elements that can generate cycles containing all 12 elements of the group, namely 6, 7, and 11. Elements whose order is equal to the order of the group

0	1	x	$x + 1$
x^2	$x^2 + 1$	$x^2 + x$	$x^2 + x + 1$
x^3	$x^3 + 1$	$x^3 + x$	$x^3 + x + 1$
$x^3 + x^2$	$x^3 + x^2 + 1$	$x^3 + x^2 + x$	$x^3 + x^2 + x + 1$

Figure 2.5

are called generators. Multiplication modulo a prime number forms a group with similar properties to those discussed considering multiplication modulo 13. The order of an element is always a divisor of the group order by Lagrange's theorem. Because some elements are generators of the group itself, it is called a cyclic group. The number of generators of the "group multiplication modulo a prime number" p having the full order $p - 1$, is $\varphi(p - 1)$ where $\varphi(\cdot)$ is Euler's totient function, namely:

$$\varphi(n) = n \prod_{p|n} \left(1 - \frac{1}{p}\right)$$

It counts the positive integers up to a given integer n that are relatively prime to n . For $p = 13$ in the previous example, we have that $\varphi(12) = 12 * \frac{1}{2} * \frac{2}{3} = 4$ which corresponds to what we have observed, i.e. that the 4 elements 2, 6, 7, and 11 are generators.

Consider again a generator $g = 2$ of the group multiplication modulo the prime $p = 13$. We can write g^0 to denote 1, g^1 to denote 2, g^2 to denote 4, this allows us to list the elements of the group as powers of the generator, i.e. $1 = g^0$, $2 = g^1$, $4 = g^2$, ... and this can be done for any group modulo a prime.

A ring is a group such that:

- the group binary operation is commutative
- has another binary operation which is: closed over the ring's elements, associative, has an identity element and distributes over the group operation

A finite field has:

- a finite set of elements
- two binary operations, which are abstract analogues to addition and multiplication
- analogues to subtraction and division using additive and multiplicative inverses

A **field** is a ring where both operations are commutative and have inverses. Intuitively a ring has addition, subtraction, and multiplication well-defined, a field adds division. A finite field is also called Galois field and integer arithmetic modulo a prime is a finite field.

The simplest Galois field is $GF(2)$: it contains two elements (0 and 1) and the binary operations are addition and multiplication both modulo 2. On top of $GF(2)$ it is possible to construct: $GF(2)[x]$ that is the set of polynomials in x with only the elements of $GF(2)$ allowed as coefficients and $GF(2)[x]/p$ that is the quotient of the set of polynomials by p .

Consider the polynomial $p(x) = x^4 + x + 1$. In this case, $GF(2)[x]/p$ contains all possible polynomials of degree less than 4, namely 2.5

Bits (n)	Feedback polynomial	Taps	Taps (hex)	Period ($2^n - 1$)
2	$x^2 + x + 1$	11	0x3	3
3	$x^3 + x^2 + 1$	110	0x6	7
4	$x^4 + x^3 + 1$	1100	0xC	15
5	$x^5 + x^3 + 1$	10100	0x14	31
6	$x^6 + x^5 + 1$	110000	0x30	63
7	$x^7 + x^6 + 1$	1100000	0x60	127
8	$x^8 + x^6 + x^5 + x^4 + 1$	10111000	0xB8	255
9	$x^9 + x^5 + 1$	100010000	0x110	511
10	$x^{10} + x^7 + 1$	1001000000	0x240	1,023
11	$x^{11} + x^9 + 1$	10100000000	0x500	2,047
12	$x^{12} + x^{11} + x^{10} + x^4 + 1$	111000001000	0xE08	4,095
13	$x^{13} + x^{12} + x^{11} + x^8 + 1$	1110010000000	0x1C80	8,191
14	$x^{14} + x^{13} + x^{12} + x^2 + 1$	11100000000010	0x3802	16,383
15	$x^{15} + x^{14} + 1$	110000000000000	0x6000	32,767
16	$x^{16} + x^{15} + x^{13} + x^4 + 1$	1101000000001000	0xD008	65,535
17	$x^{17} + x^{14} + 1$	10010000000000000	0x12000	131,071
18	$x^{18} + x^{11} + 1$	100000010000000000	0x20400	262,143
19	$x^{19} + x^{18} + x^{17} + x^{14} + 1$	1110010000000000000	0x72000	524,287
20	$x^{20} + x^{17} + 1$	10010000000000000000	0x90000	1,048,575
21	$x^{21} + x^{19} + 1$	10100000000000000000	0x140000	2,097,151
22	$x^{22} + x^{21} + 1$	110000000000000000000	0x300000	4,194,303
23	$x^{23} + x^{18} + 1$	1000010000000000000000	0x420000	8,388,607
24	$x^{24} + x^{23} + x^{22} + x^{17} + 1$	111000010000000000000000	0xE10000	16,777,215

Figure 2.6

Polynomials $p(x)$ such that x is a generator of $GF(2)[x]/p(x)$ are called primitive polynomials. In practice, this means that it is possible to start with 1 and, by using the procedure below, it is possible to generate all the elements of $GF(2)[x]/p(x)$

- consider 1
- multiply by x and if the result contains a term x^N with $N = \text{degree of } p(x)$, then subtract $p(x)$
- stop when the result is 1, this happens when the power of the generator is $2^N - 1$

Primitive polynomials $p(x)$ allow us to design maximal LFSRs by exploiting the following correspondence between these and $GF(2)[x]/p(x)$. In conclusion $GF(2)[x]/p(x)$ "corresponds" to $GF(2^N)$ when $p(x)$ is a primitive polynomial of degree N . These are useful because guarantee to generate the longest possible cycle.

Given a maximum-length LFSR of n bits and reading $2^n - 1$ consecutive bits of the m-sequence that it produces, we have that:

- one half of the bits are 1 and one half are 0 (actually the 1's are one more than the 0's)

- there are 2^{n-1} runs: $1/2$ of the runs has length 1, $1/4$ of the runs has length 2, ..., $1/2^i$ of the runs has length i (for $2 \leq i \leq n-2$), there is only one run of $n-1$ zeros and none of the runs has $n-1$ ones, there is only one run of n ones and none of the runs has n zeros.

As an example consider using as feedback polynomial the primitive polynomial $x^5 + x^2 + 1$, hence the period is $2^5 - 1 = 31$ bits and we expect $2^{5-1} = 16$ runs. Starting from the state (10000), the obtained m-sequence is:

$$0000 - 1 - 00 - 1 - 0 - 11 - 00 - 11111 - 000 - 11 - 0 - 111 - 0 - 1 - 0 - 1$$

There are indeed 16 runs and 8 runs have length 1, 4 runs have length 2, 2 runs have length 3, there is only 1 run of 4 zeros and none of 4 ones and there is only 1 run of 5 ones and none of 5 zeros.

2.8 LFSR and linear algebra

For analysing LFSRs, it is useful to investigate their relationships with Linear Algebra. For this, the key idea is the notion of companion matrix: given the polynomial $p(x) = c + c_1 * x + \dots + c_n * x^n$ in $GF(2)[x]$ its companion matrix is:

$$C(p) = \begin{bmatrix} 0 & 0 & \dots & 0 & c_0 \\ 1 & 0 & \dots & 0 & c_1 \\ 0 & 1 & \dots & 0 & c_2 \\ \vdots & \vdots & \ddots & \vdots & \vdots \\ 0 & 0 & \dots & 1 & c_{n-1} \end{bmatrix}$$

It is possible to show that multiplication by C is equivalent to multiplication by x in $GF(2)[x]/p(x)$. This connection allows us to consider the state recovery problem, i.e. predicting the state vector of the LFSR from its output bits. If we can solve this we can recover the initial state or, equivalently, the seed. And this requires polynomial time computations, this leads to the LFSRs main weaknesses: LFSRs can not be used directly in cryptography because of their linearity. Given any LFSR of length n , the j -th bit of the keystream, for can be obtained as a linear combination of its previous n bits keystream.

Let the bits of the output sequence denoted by a_0, \dots, a_n , there exists n bits $\lambda_0, \dots, \lambda_n$ such that $\sum_{i=0}^n \lambda_i a_i = 0$, $\lambda_n = 1$, note that the values $\lambda_0, \dots, \lambda_n$ do not depend on when Eve starts intercepting the ciphertext.

Knowing the values of $\lambda_0, \dots, \lambda_n$ is equivalent to knowing the feedback polynomial: $g = \sum_{i=0}^n \lambda_i x^i$ and thus being able to predict the keystream. Eve may obtain the required subsequence of the keystream mounting a known-plaintext or chosen-plaintext attack, hence LFSRs are vulnerable with respect to known-plaintext attacks and they must not be used as keystream generators, although they can be used as components in particular when iterated and combined by using some kind of non-linear functions. Linearity is good because it is easy to implement, but on the other hand is not secure enough.

2.8.1 DVD encryption

DVD encryption uses the CSS stream cipher to encrypt movie contents using a 40-bit secret key. The CSS stream cipher is particularly weak as it can be broken in far less time than an exhaustive search over all 2^{40} seeds 2.7.

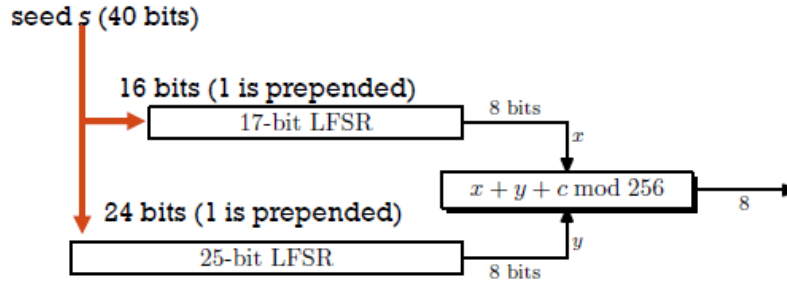


Figure 2.7: DVD CSS Encryption

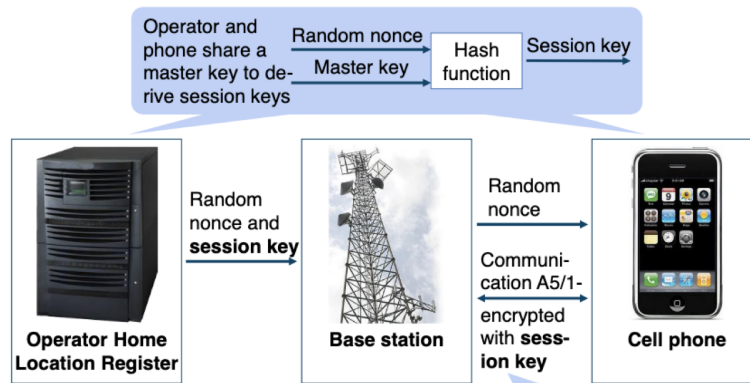


Figure 2.8: GSM

The two LFSRs are run in parallel for 8 cycles and then the resulting bits are considered as integers and added modulo 256. The simplest attack is simple brute-force: suppose Eve knows the first 100 bytes of the output sequence, then just guess a 40 bits seed, run for 100 iterations and compare with the output sequence, if they match the seed was found, otherwise try another seed. And even smarter approach could be used to only attempt at most 2^{16} seeds.

2.9 A5 family of ciphers

2.9.1 GSM

The GSM (Global System for Mobile communication) standard was designed from 1982-1991, the level of security specifications regarding both authentication and encryption were limited: algorithms were never officially published (achieving security by obscurity), thus algorithms were reverse-engineered or leaked, leading to revelations of several possible attacks, and within a few months after the release, most of the cryptographic schemes had been compromised and some were even proven to be close to useless 2.8.

Each frame is numbered by a frame number, obtained by a frame counter initialized with 0 at conversation-start and incremented by $1 \bmod 2^{22}$ with each frame sent. An algorithm of the A5 family takes the session key K_c (symmetric) and a frame counter Fn and generates 228 pseudo random bits (PRAND) called a key stream. The keystream is then XORed with a 228 bit segment

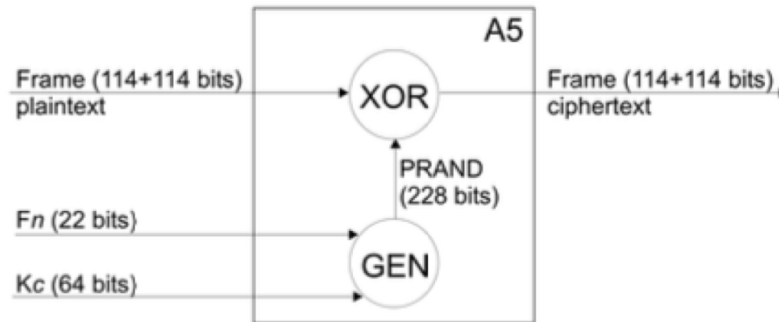


Figure 2.9: A5 ciphers operation




Figure 2.10: A5/1

of plain text yielding 228 bits of ciphertext 2.9.

2.9.2 A5/0

It is the weakest of the A5 versions as it offers no encryption. It is a no-operation cipher, that generates the pseudo random bits by negating the input frame, thus leaving out the XOR function. The result is an algorithm that outputs the plain text it received as an input. This version is found in third world countries or countries with UN sanctions.

2.9.3 A5/1

It uses 3 LFSRs (with lengths 19, 22, 23, total 64 bits of state) and has a keystream length of 228 bits 2.10.

When a register is clocked, its tapped bits are XORed and the result is stored in the registers LSB, the register MSB is shifted out of the register, and its value is forgotten.

In short first a seed (secret vector K of 64 bits) is selected from the session key, then an initialization vector (public vector IV of 22 bits) that is the frame counter, from K and IV we get the initial state (namely a vector of 64 bits) using a keyloading function $kl(K, IV)$. The function kl is defined by iteration: initially the registers contain all zeros, the bits in K are injected by xor-ing them with other bits, and similarly for the bits of IV . After the keyloading the warmup starts, the registers R1, R2, and R3 are irregularly clocked 100 times, without producing output, and it is as follows: in each step, the clock bits, each one among R1, R2, and R3 is updated if the clock bits

agree with the majority of the clock bits. In this way, it is possible to show that each register clocks with a probability of $3/4$. At this point, the initialization of the registers is complete and the stream cipher is ready to output the key-stream.

The keystream is composed of 228 bits that are obtained in 228 steps, at each clock tick the update function is defined as follows:

- The register R1 is updated if its clock bit agrees with the majority of the others
- The register R2 is updated if its clock bit agrees with the majority of the others
- The register R3 is updated if its clock bit agrees with the majority of the others

The output function is defined as: the xor of the output of the three registers R1, R2, R3 is the output bit. After 228 bits of output, the key-loading phase is executed again.

2.9.4 A5/2

The key stream length remains the same as A5/1 but the number of LFSRs is increased to 4 respectively with lengths 19,22,23,17 with a total 81 bits of state.

In short a seed (secret vector K of 64 bits) is selected, and an initialization vector (public vector IV of 22 bits), from K and IV we get the initial state (namely a vector of 64 bits) using a keyloading function $kl(K, IV)$. kl and warmup are defined similarly to A5/1.

At each clock tick the update function is defined as follows:

- R_1 is updated if $R_4[6] = \text{maj}(R_4[6], R_4[13], R_4[9])$
- R_2 is updated if $R_4[13] = \text{maj}(R_4[6], R_4[13], R_4[9])$
- R_3 is updated if $R_4[9] = \text{maj}(R_4[6], R_4[13], R_4[9])$

For the output a majority function is applied as non-linear filtering function to each of the first three registers. After a month an attack was found that could break the cipher almost in real time with a known plaintext attack.

2.9.5 A5/3

A5/3 is the last stream cipher of the A5 family and provides users with a higher level of security than both A5/1 and A5/2. It is based on the block cipher KASUMI. The key space is 128 bits and the message space is 64 bits 2.11.

The problem is that the key generated Kc is only 64 bits, so the maximal exhaustive search complexity is (only) 2^{64} . To make things worse Kc is generated only once after the cell phone registers with the network and stays active for all communication, until the telco requests a new one or the cell phone deregisters (so if one key is compromised every other communication is compromised until a new key is requested), and Kc is artificially shortened in deployed systems when zeroing 10 bits, lowering search complexity to 2^{54} , and encryption is applied after error correction.

2.10 Remarks on the A5 Family

- A5/1 is affected by a number of serious weaknesses, and its use is strongly discouraged, since there are practical attacks that can break the cipher

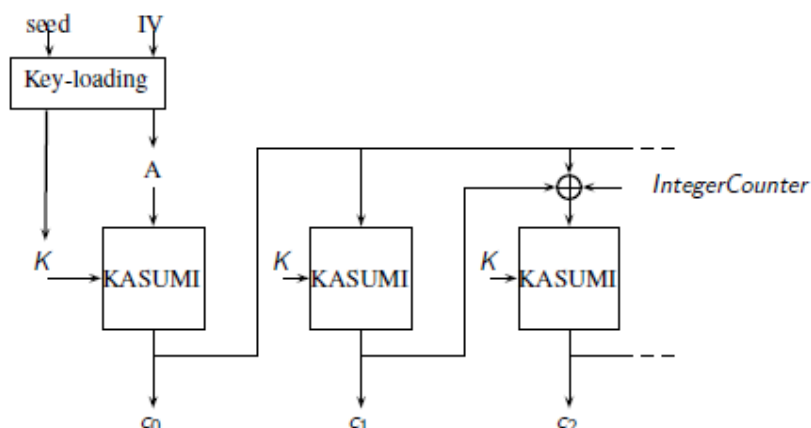


Figure 2.11

- A5/2 is extremely weak and it can be broken in real time with inexpensive equipment; it is therefore no longer supported by new mobile phones
- A5/3 is the common standard for the new generation of mobile and it is considered secure, even though there do exist practical attacks to KASUMI that suggest some significant weaknesses of the cipher

2.11 Practicality of attacking GSM communication

In theory, the attacks are relatively simple. In practice, a considerable amount of hardware is necessary to actually intercept GSM communications. The hardware must at least consist of a radio receiver device which is capable of receiving and decoding digital data that is exchanged over-the-air. Hypothetically a simple GSM mobile phone already has all these capabilities (except the decrypting of an unknown A5 stream), so it might be possible to use such a phone for eavesdropping nevertheless a huge amount of know-how, time and money is needed.

2.12 Bluetooth

Bluetooth is a wireless technology standard invented by Ericsson in 1994 for exchanging data over short distances using short-wavelength UHF radio waves (Range: 2.4 to 2.485 GHz) from fixed and mobile devices. The standard offers methods for generating keys, authenticating users, and encrypting data. BLE was introduced in 2011 as Bluetooth 4.0, the main difference between BLE and Bluetooth is power consumption. BLE uses the AES cipher with 128-bit key length to provide data encryption and integrity over the wireless link.

2.12.1 E0

Used in Bluetooth standard and introduced in 1999. It uses 4 linear registers with lengths 25,31,33,39 with a total linear par of 128 bits. It has also one non linear register that is 4 bits. The update functions depend on the non-linear register 2.12.

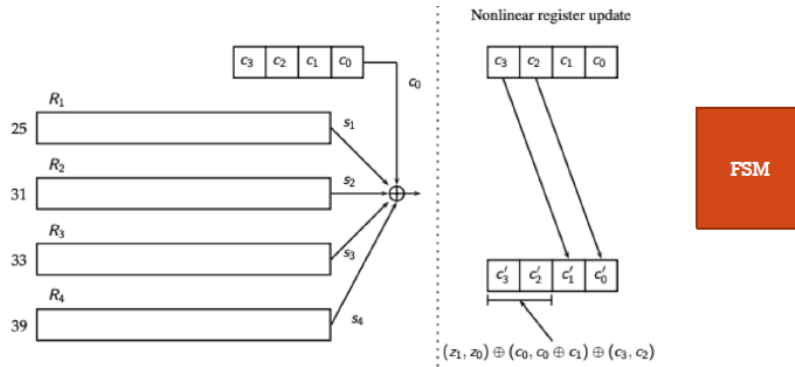


Figure 2.12

2.13 RC4

Rivest Cipher 4 was designed in 1987. At first its implementation was trade secret, but then someone reverse engineered it, later Rivest confirmed that the code that was leaked was in fact correct. RC4 is used mainly because it is easy to implement in hardware and it is fast. RC4 was rapidly adopted in commonly used encryption protocols and standard like WEP and SSL.

RC4 generates a pseudorandom stream of bytes, the key stream. To generate the keystream, the cipher makes use of a secret internal state which consists of two parts: a permutation of all 256 possible bytes (denoted S) and two 8-bit index-pointers (denoted i and j). The permutation is initialized with a variable length key, typically between 40 and 2048 bits, using the so-called key-scheduling algorithm (KSA), once this has been completed, the stream of bits is generated using the pseudorandom generation algorithm (PRGA). While many stream ciphers are based on LFSRs (efficient in hardware but less so in software), the design of RC4 avoids the use of LFSRs and is ideal for software implementation, as it requires only byte manipulations.

The basic data structure needed is an array of 256 8-bit integers, called the state vector S that is initialized with the encryption key T with the following procedure:

1. S is initialized with entries from 0 to 255 in the ascending order
2. S is further initialized with the help of a temporary 256-element vector denoted T that also holds 256 integers. The vector T is initialized with the encryption key as follows: let K be the encryption key represented as a vector 8-bit integers of size 16 (in case of a 128-bit key), i.e. K stores 16 non-negative integers whose values will be between 0 and 255, then initialize the 256-element vector T by placing in it as many repetitions of the key as necessary until T is full
3. Use the 256-element vector T to produce the initial permutation of S as in 2.13

The keystream is generated with the algorithm in 2.14

But why use $j = (j + S[i] + T[i]) \bmod 256$? Because it works in practice. It is also easy to compute. But also the way you pick indexes is biased and this introduces vulnerabilities.

Theoretical analysis shows that for a 128 bit key length, the period of the pseudorandom sequence of bytes is likely to be greater than 10^{100} , but it has been shown to be vulnerable to attacks especially if the beginning portion of the output pseudorandom byte stream is not discarded. As a consequence use of RC4 is prohibited in SSL/TLS protocol since 2015.


```

j = 0
for i = 0 to 255
    j = ( j + S[i] + T[i] ) mod 256
    SWAP S[i], S[j]

```

Figure 2.13: Key Scheduling Algorithm (KSA)

```

i, j = 0
while ( true )
    i = ( i + 1 ) mod 256
    j = ( j + S[i] ) mod 256
    SWAP S[i], S[j]
    k = ( S[i] + S[j] ) mod 256
    output S[k]

```

Figure 2.14: Generating the pseudorandom byte stream

WiFi security started with RC4 in the WEP protocol. After it was discovered that the encryption key used in WEP could be acquired by an adversary in almost no time, WiFi security has now moved on to the WPA2 protocol that uses AES for encryption.

Unlike modern stream ciphers, RC4 does not take a separate nonce alongside the key. If a single long-term key is to be used to securely encrypt multiple streams, the protocol must specify how to combine the nonce and the long-term key to generate the stream key for RC4. One approach to addressing this is to generate a "fresh" RC4 key by hashing a long-term key with a nonce. Unfortunately, many applications that use RC4 simply concatenate key and nonce; this gives rise to related key attacks that are famous for breaking the WEP standard. Because RC4 is a stream cipher, it is more malleable than block ciphers. If not used together with other techniques (e.g., message authentication codes), then encryption is vulnerable to bit-flipping attacks.

Stream ciphers require a shared secret, a seed, which should be transmitted via a secure channel and once Alice and Bob have the seed, they can exchange via an insecure channel other parameters, like the IV and start encrypting/decrypting. The use of stream ciphers guarantees confidentiality. However, the stream cipher in itself does not guarantee that the seed has been exchanged securely, this has to be done in other ways. On the other hand, the attacker might tamper with the ciphertext and Alice/Bob will not understand immediately that the messages have been corrupted. In other words, stream ciphers provide neither authentication nor integrity.

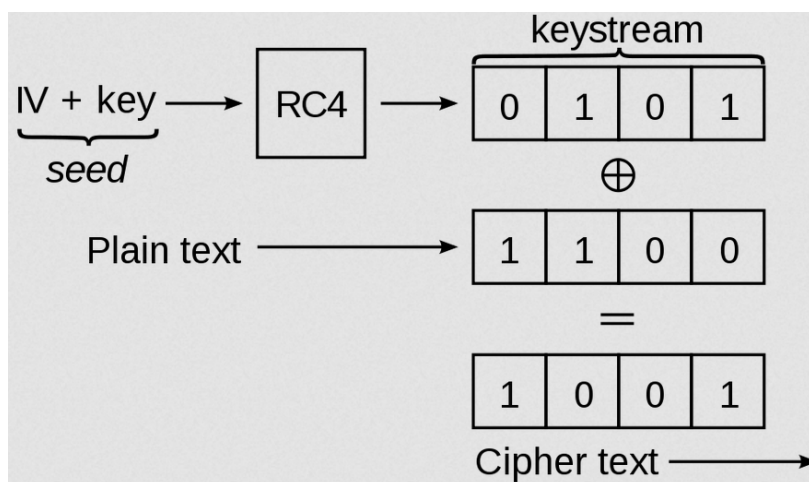


Figure 2.15

2.14 RC4 in WEP

2.14.1 WIFI Protocol

- Wired Equivalent Privacy (WEP) is the oldest protocol and has been proven to be vulnerable.
- Wi-Fi Protected Access (WPA) improved security but still vulnerable.
- Wi-Fi Protected Access II (WPA2) while not perfect is the most secure choice. It used two different types of cryptographic techniques to secure communication and those are Temporal Key Integrity Protocol (TKIP) and Advanced Encryption Standard (AES).

TKIP is used to authenticate the client and exchange messages in a secure way. Authentication is based off the Pre Shared Key, a hard wired string written on the device.

2.14.2 WIFI Authentication

Client and the station have to agree on a key used to encrypt messages using symmetric encryption. How do we share this secret key? They have to agree over an insecure channel, so we need a protocol to do this. So the Client reads this preshared key on the device to authenticate to the device, but this key won't be used to encrypt messages! Because it is shared between clients! So TKIP used a key mixing function that takes as input this preshared key and an initialization vector and passing it to RC4 cipher initialization 2.15.

2.14.3 Security of WPA2

Vulnerable to KRACK attack. It exploits a functionality where if a user is disconnected from WiFi then for reconnecting it performs an easier problem to derive a new session.